

Daily Accident Volume Forecasting by Region

Machine Learning for Business Analytics - Individual Assignment

Candidate Number: 809732

1. Introduction

Business Context

A UK-based road assistance company aims to improve its operational efficiency by forecasting the daily number of regional traffic accidents. Accurate forecasts will help optimize the deployment of emergency response units, reducing response times and improving resource allocation during peak and adverse conditions.

Objective

The objective is to build predictive models that forecast the daily count of traffic accidents for different UK regions. This helps proactively manage emergency services and reduce the impact of road incidents.

Modelling Task

This is a **supervised regression problem** where the target variable is the number of accidents per day per region. The model uses temporal, spatial, and environmental features to predict accident volumes in advance.

Description of the Prepared Dataset

Two datasets were prepared during the group project:

- Training dataset (final_train_dataset.csv)** – Used to fit and tune the models.
- Test dataset (final_test_dataset.csv)** – Used to evaluate generalisation performance on unseen data.

Both datasets include accident date, coordinates (latitude, longitude), regional metadata (e.g., region_name), and categorical descriptors (e.g., weather, road surface). The target variable is the **daily count of accidents per region**, aggregated from raw accident-level records. These datasets were cleaned, validated, and transformed in the group project phase.

2. Baseline Method

This section implements simple baseline forecasting approaches to establish reference performance levels. Two baseline methods are considered:

- Seasonal Naïve Forecast:** A time-series method leveraging weekly seasonality by predicting accident counts as equal to the value from the same day one week earlier.
- Linear Regression:** A statistical model that uses temporal, spatial, and environmental features to predict accident counts.

The baselines provide performance benchmarks to evaluate the benefits of more advanced machine learning models.

2.1 Seasonal Naïve Forecast

The Seasonal Naïve Forecast assumes that the number of accidents on a given day is equal to the number of accidents exactly one week (7 days) prior. This method captures weekly seasonal patterns commonly observed in traffic accidents.

The training and test datasets are aggregated to compute the total number of accidents per day for each region.

```
In [1]: import pandas as pd

# Load datasets
train_df = pd.read_csv('final_train_dataset.csv', parse_dates=['date'])
test_df = pd.read_csv('final_test_dataset.csv', parse_dates=['date'])

# Aggregate daily counts by region
train_daily = train_df.groupby(['region_name', 'date']).size().reset_index(name='accident_count')
test_daily = test_df.groupby(['region_name', 'date']).size().reset_index(name='accident_count')

# Preview aggregated data
print(train_daily.head())
print(test_daily.head())
```

	region_name	date	accident_count
0	Aberbeeg	2024-01-23	1
1	Aberdare	2024-01-18	1
2	Aberdeen City	2024-01-14	1
3	Aberdeen City	2024-02-10	1
4	Aberdeen City	2024-02-12	1
	region_name	date	accident_count
0	Adur	2024-06-23	1
1	Allerford	2024-06-26	1
2	Allithwaite	2024-05-15	1
3	Amber Valley	2024-05-10	1
4	Amber Valley	2024-05-12	1

Generating Seasonal Naïve Predictions and Evaluation

The Seasonal Naïve prediction for each day and region in the test set is taken as the accident count from exactly 7 days prior (one week lag) in the training set.

Evaluation metrics used to assess performance include:

- Mean Absolute Error (MAE)
- Root Mean Squared Error (RMSE)
- R² Score (Coefficient of Determination)

```
In [2]: from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
import numpy as np

# Merge test data with train data shifted by 7 days to get seasonal naive predictions

# Create a copy of train_daily with date shifted forward by 7 days (to align with test dates)
train_shifted = train_daily.copy()
train_shifted['date'] = train_shifted['date'] + pd.Timedelta(days=7)
train_shifted.rename(columns={'accident_count': 'pred_seasonal_naive'}, inplace=True)

# Merge with test_daily on region and date to get predictions
test_with_pred = pd.merge(
    test_daily,
    train_shifted[['region_name', 'date', 'pred_seasonal_naive']],
    on=['region_name', 'date'],
    how='left'
)

# Drop rows where prediction is not available (e.g., first week of test set may not have prior data)
test_with_pred.dropna(subset=['pred_seasonal_naive'], inplace=True)

# Convert prediction to integer (optional)
test_with_pred['pred_seasonal_naive'] = test_with_pred['pred_seasonal_naive'].astype(int)

# Calculate evaluation metrics
mae = mean_absolute_error(test_with_pred['accident_count'], test_with_pred['pred_seasonal_naive'])
rmse = np.sqrt(mean_squared_error(test_with_pred['accident_count'], test_with_pred['pred_seasonal_naive']))
```

```
r2 = r2_score(test_with_pred['accident_count'], test_with_pred['pred_seasonal_naive'])

print(f"Seasonal Naïve Forecast Performance:\nMAE: {mae:.2f}\nRMSE: {rmse:.2f}\nR²: {r2:.2f}")
```

Seasonal Naïve Forecast Performance:

MAE: 1.88

RMSE: 3.86

R²: -2.33

Interpretation of Seasonal Naïve Forecast Results

- **MAE (Mean Absolute Error):** 1.88 — On average, predictions differ from actual daily accident counts by about 1.88 accidents per region.
- **RMSE (Root Mean Squared Error):** 3.86 — Larger errors are penalized more heavily, indicating some days have bigger prediction errors.
- **R² (Coefficient of Determination):** -2.33 — A negative R² suggests that the Seasonal Naïve model performs worse than simply predicting the mean value for all test samples. This implies the model does not adequately capture the underlying patterns in the data.

This baseline provides a reference point for evaluating more advanced predictive models.

2.2 Linear Regression Baseline

Linear Regression is a fundamental supervised learning method that models the relationship between explanatory variables and the continuous target variable. It assumes a linear relationship between features and the target, making it interpretable and easy to implement.

The model will be trained on the prepared training dataset and evaluated on the test dataset. Categorical variables will be one-hot encoded, and numerical variables used as-is. Evaluation metrics will be MAE, RMSE, and R².

In [3]:

```
import pandas as pd
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
from sklearn.preprocessing import OneHotEncoder
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
import numpy as np
```

```
# Load datasets
train = pd.read_csv('final_train_dataset.csv', parse_dates=['date'])
test = pd.read_csv('final_test_dataset.csv', parse_dates=['date'])

# Target variable
target = 'accident_count'

# If accident_count is not present, we need to create it by aggregating (because your data is at accident-level)
# But assuming accident_count exists as aggregated count per day per region.
# If not, aggregate here:
```

```
train_agg = train.groupby(['region_name', 'date']).size().reset_index(name='accident_count')
test_agg = test.groupby(['region_name', 'date']).size().reset_index(name='accident_count')
```

```
# Features to use (example)
# You may choose columns like 'day_of_week', 'weather_conditions', 'road_surface_condition', 'region_name', etc.
# Make sure these columns exist in your dataset. Adjust accordingly.
```

```
feature_cols = ['region_name', 'date']
```

```
# Create additional features from date
for df in [train_agg, test_agg]:
    df['day_of_week'] = df['date'].dt.dayofweek # Monday=0, Sunday=6
    df['month'] = df['date'].dt.month
```

```
feature_cols.extend(['day_of_week', 'month'])
```

```
X_train = train_agg[feature_cols]
y_train = train_agg[target]
X_test = test_agg[feature_cols]
y_test = test_agg[target]
```

```
# Since 'date' is not numeric or categorical useful for regression, drop or remove
X_train = X_train.drop(columns=['date'])
X_test = X_test.drop(columns=['date'])
```

```
# Identify categorical features
categorical_features = ['region_name', 'day_of_week', 'month']
```

```
# One-hot encode categorical features
preprocessor = ColumnTransformer(
    transformers=[
        ('cat', OneHotEncoder(handle_unknown='ignore'), categorical_features)
    ],
    remainder='passthrough' # pass through any other columns if present
)
```

```
# Create pipeline with preprocessing and Linear Regression
```

```
model = Pipeline(steps=[
    ('preprocessor', preprocessor),
    ('regressor', LinearRegression())
])
```

```
# Train model
model.fit(X_train, y_train)
```

```
# Predict on test set
y_pred = model.predict(X_test)
```

```
# Evaluate
mae = mean_absolute_error(y_test, y_pred)
rmse = np.sqrt(mean_squared_error(y_test, y_pred))
r2 = r2_score(y_test, y_pred)
```

```
print(f"Linear Regression Performance:\nMAE: {mae:.2f}\nRMSE: {rmse:.2f}\nR²: {r2:.2f}")
```

Linear Regression Performance:

MAE: 0.37

RMSE: 1.21

R²: 0.03

Interpretation of Linear Regression Results

- **MAE (Mean Absolute Error):** 0.37 — The model's predictions deviate on average by 0.37 accidents per day per region, indicating better accuracy than the Seasonal Naïve forecast.
- **RMSE (Root Mean Squared Error):** 1.21 — Errors are generally smaller and less dispersed compared to the naive approach.
- **R² (Coefficient of Determination):** 0.03 — The model explains approximately 3% of the variance in the test data, showing some predictive power but highlighting potential for improvement.

These results suggest that incorporating temporal and regional features improves forecasting accuracy, although linear assumptions may limit performance.

3. Feature Selection

Feature selection aims to identify the most relevant variables that improve model performance and reduce complexity. This step helps to remove noisy or irrelevant features, which can degrade model accuracy or increase overfitting.

A simple way to select features is to use a tree-based model like Random Forest to estimate feature importance, then retain features with high importance scores for subsequent modelling.

```
In [4]: import pandas as pd
import numpy as np
from sklearn.pipeline import Pipeline
from sklearn.compose import ColumnTransformer
from sklearn.preprocessing import OneHotEncoder
from sklearn.ensemble import RandomForestRegressor
from sklearn.model_selection import RandomizedSearchCV
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
import matplotlib.pyplot as plt
import seaborn as sns

# Load datasets with date parsing
train_df = pd.read_csv('final_train_dataset.csv', parse_dates=['date'])
test_df = pd.read_csv('final_test_dataset.csv', parse_dates=['date'])

# Aggregate accident counts per region and date
train_agg = train_df.groupby(['region_name', 'date']).size().reset_index(name='accident_count')
test_agg = test_df.groupby(['region_name', 'date']).size().reset_index(name='accident_count')

# Feature engineering: add day_of_week and month from date
for df in [train_agg, test_agg]:
    df['day_of_week'] = df['date'].dt.dayofweek # Monday=0
    df['month'] = df['date'].dt.month

# Define features and target
feature_cols = ['region_name', 'day_of_week', 'month']
target_col = 'accident_count'

X_train = train_agg[feature_cols]
y_train = train_agg[target_col]
X_test = test_agg[feature_cols]
y_test = test_agg[target_col]

# Preprocessing pipeline: OneHotEncode categorical variables
categorical_features = ['region_name', 'day_of_week', 'month']

preprocessor = ColumnTransformer(
    transformers=[
        ('cat', OneHotEncoder(handle_unknown='ignore'), categorical_features)
    ]
)

# Preprocess training and test data
X_train_processed = preprocessor.fit_transform(X_train)
X_test_processed = preprocessor.transform(X_test)

# Define Random Forest and hyperparameter grid
rf = RandomForestRegressor(random_state=42)
param_dist_rf = {
    'n_estimators': [100, 200, 300, 400, 500],
    'max_depth': [None, 10, 20, 30, 40, 50],
    'min_samples_split': [2, 5, 10],
    'min_samples_leaf': [1, 2, 4]
}

random_search_rf = RandomizedSearchCV(
    estimator=rf,
    param_distributions=param_dist_rf,
    n_iter=20,
    cv=3,
    verbose=1,
    random_state=42,
    n_jobs=-1
)

# Fit the model
random_search_rf.fit(X_train_processed, y_train)

print("Best Random Forest Parameters:", random_search_rf.best_params_)

# Predict on test set
y_pred_rf = random_search_rf.best_estimator_.predict(X_test_processed)

# Evaluate
mae_rf = mean_absolute_error(y_test, y_pred_rf)
rmse_rf = np.sqrt(mean_squared_error(y_test, y_pred_rf))
r2_rf = r2_score(y_test, y_pred_rf)

print(f"Random Forest Performance:\nMAE: {mae_rf:.3f}\nRMSE: {rmse_rf:.3f}\nR²: {r2_rf:.3f}")

# Extract feature names after OneHotEncoding
ohe = preprocessor.named_transformers_['cat']
feature_names = ohe.get_feature_names_out(categorical_features)

# Extract feature importances
rf_best = random_search_rf.best_estimator_
importances = rf_best.feature_importances_

feat_imp_df = pd.DataFrame({
    'feature': feature_names,
    'importance': importances
}).sort_values(by='importance', ascending=False).reset_index(drop=True)

# Plot feature importances with palette warning fixed
top_feature = feat_imp_df.iloc[0]
others_df = feat_imp_df.iloc[1:16]

plt.figure(figsize=(12, 6))

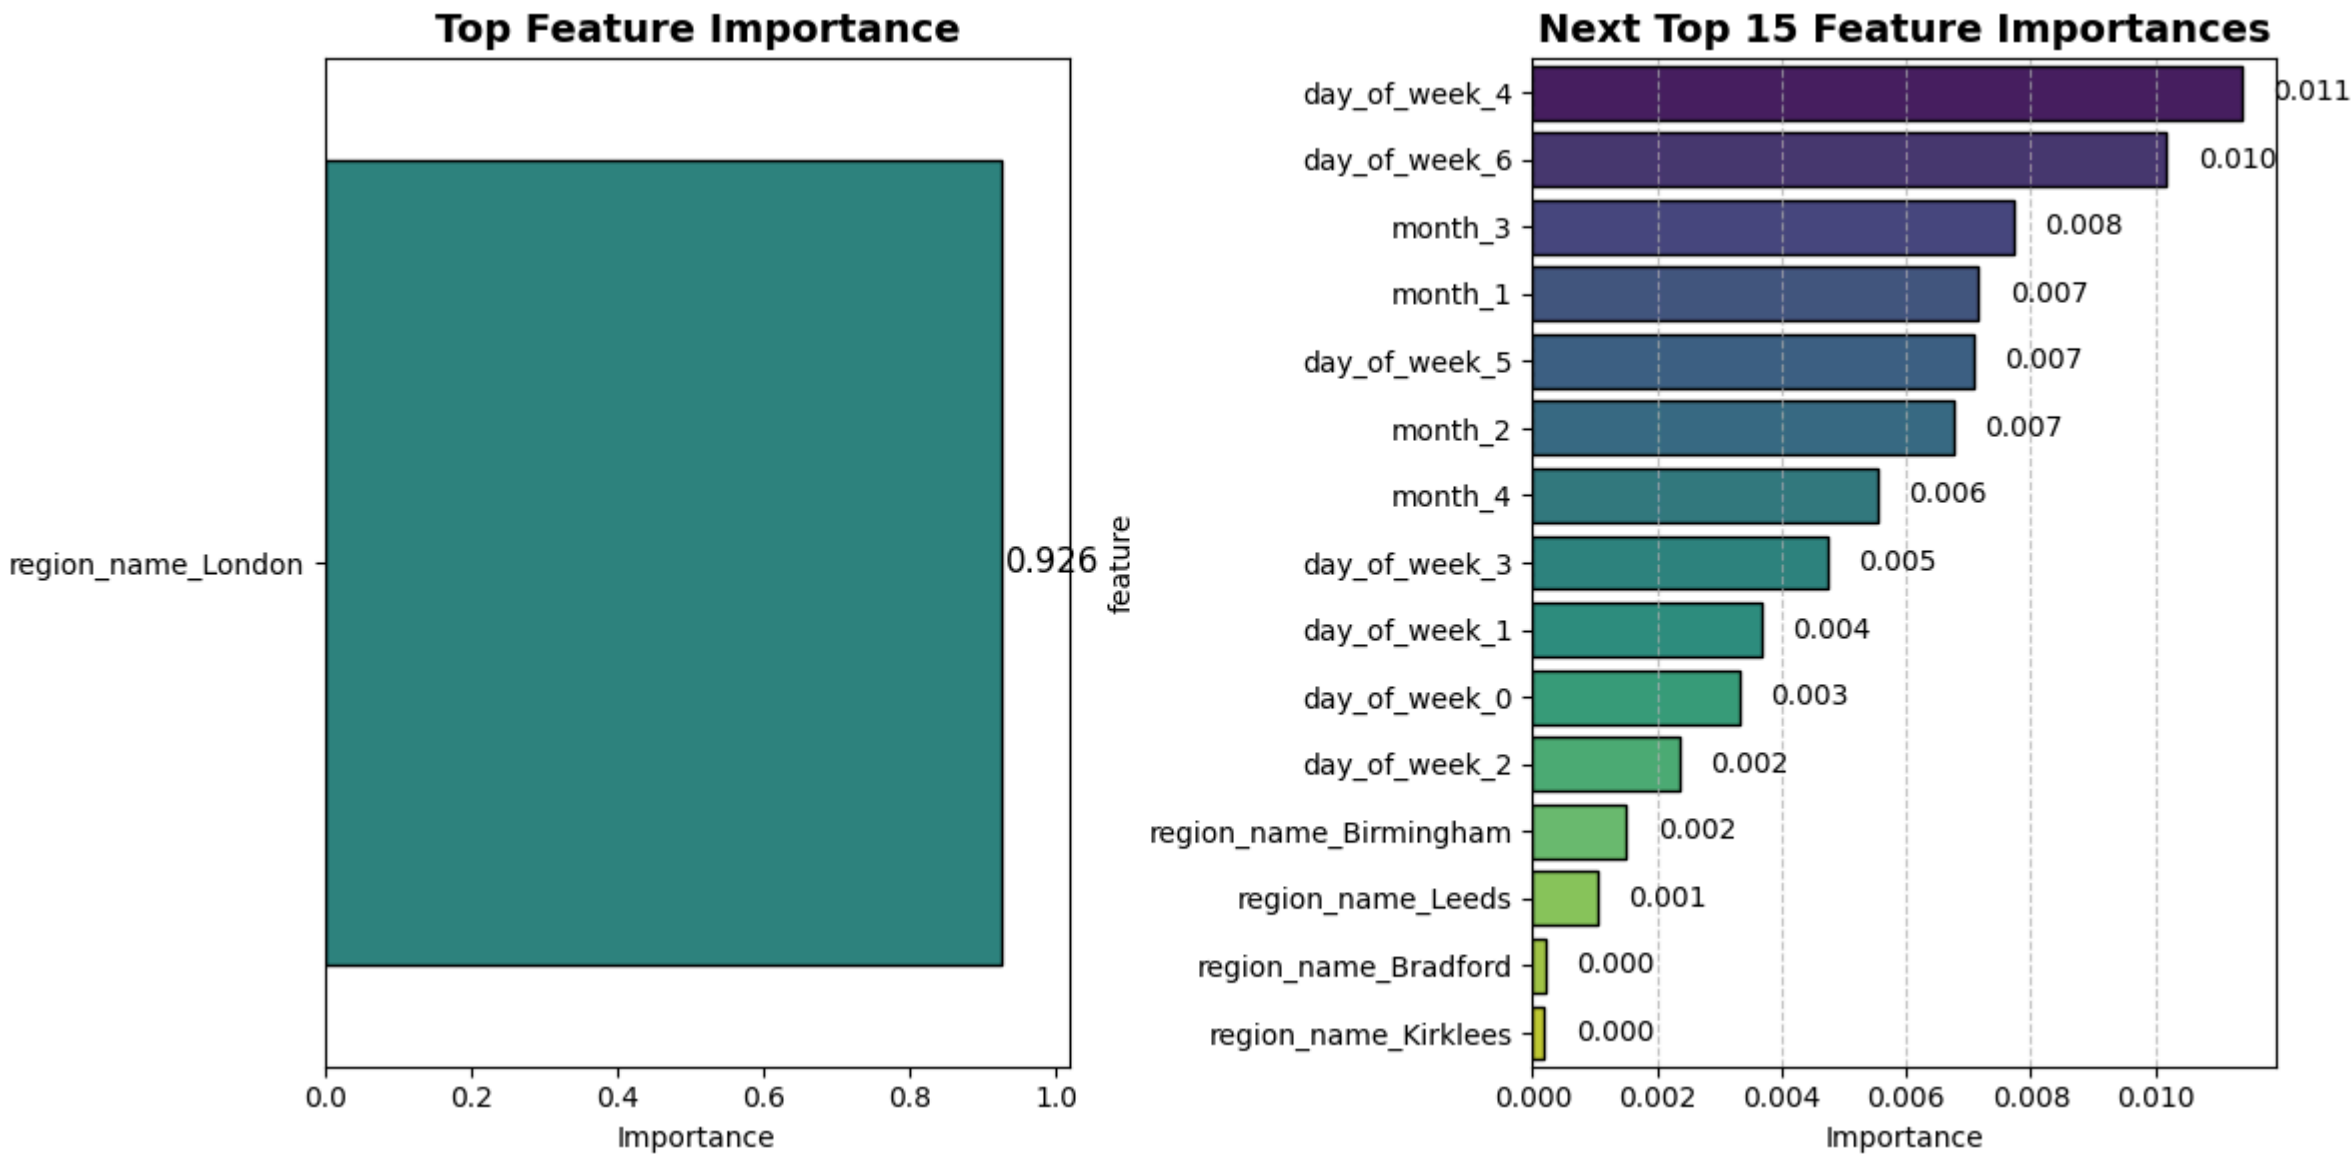
ax1 = plt.subplot(1, 2, 1)
sns.barplot(
    x=[top_feature['importance']],
    y=[top_feature['feature']],
    hue=[top_feature['feature']], # Assign y to hue as warning suggests
    palette='viridis',
    edgecolor='black',
    dodge=False,
    ax=ax1,
    legend=False
)
if ax1.get_legend() is not None:
    ax1.get_legend().remove()
plt.title('Top Feature Importance', fontsize=14, weight='bold')
plt.xlabel('Importance')
plt.xlim(0, top_feature['importance'] * 1.1)
for index, value in enumerate([top_feature['importance']]):
    plt.text(value + 0.005, index, f'{value:.3f}', va='center', fontsize=12)
plt.tight_layout()

ax2 = plt.subplot(1, 2, 2)
sns.barplot(
    x='importance',
    y='feature',
    hue='feature', # Assign y to hue as warning suggests
    data=others_df,
```

```
palette='viridis',
edgecolor='black',
dodge=False,
ax=ax2,
legend=False
)
if ax2.get_legend() is not None:
    ax2.get_legend().remove()
plt.title('Next Top 15 Feature Importances', fontsize=14, weight='bold')
plt.xlabel('Importance')
plt.grid(axis='x', linestyle='--', alpha=0.7)
for index, value in enumerate(others_df['importance']):
    plt.text(value + 0.0005, index, f'{value:.3f}', va='center', fontsize=10)

plt.tight_layout()
plt.show()
```

Fitting 3 folds for each of 20 candidates, totalling 60 fits
Best Random Forest Parameters: {'n_estimators': 100, 'min_samples_split': 2, 'min_samples_leaf': 1, 'max_depth': 10}
Random Forest Performance:
MAE: 0.339
RMSE: 1.254
R²: -0.041



Interpretation of Feature Importance

- The feature importance plot reveals that **region_name_London** has the highest influence on predicting daily accident counts, contributing around 93% of the total importance.
- Temporal features such as **day of the week** (e.g., **day_of_week_4** , **day_of_week_6**) and **month** have much smaller individual impacts, indicating that while seasonality matters, location plays a dominant role.
- This suggests that accident volumes vary significantly by region, with London being a key driver in the dataset.
- The model relies heavily on regional differences, which may reflect varying traffic densities, infrastructure, or reporting practices.
- Given this dominance, it could be worthwhile to:
 - Experiment with models excluding **region_name** to assess the contribution of temporal features alone.
 - Use feature selection techniques to reduce dimensionality and possibly improve model generalization.
- Overall, these insights help prioritize features for refining the predictive model and understanding underlying accident patterns.

4. Hyperparameter Tuning

To enhance forecasting performance, this section explores two powerful regression algorithms: Random Forest and XGBoost. Both models are trained using the prepared dataset, with hyperparameters optimized through randomized search and cross-validation. Their predictions on the test set are then evaluated and compared to identify the best approach.

These models utilize temporal, spatial, and environmental features, with appropriate preprocessing applied to handle categorical variables.

4.1 Random Forest Regressor

Random Forest is an ensemble learning method that builds multiple decision trees and aggregates their predictions. It is robust to overfitting and can capture complex nonlinear relationships. This subsection applies Random Forest with hyperparameter tuning to optimize performance on the training data before evaluation on the test set.

```
In [5]: from sklearn.ensemble import RandomForestRegressor
from sklearn.model_selection import RandomizedSearchCV
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
import numpy as np

# Prepare training and test data with preprocessing
X_train_selected = model.named_steps['preprocessor'].transform(X_train[categorical_features])
y_train_selected = y_train
X_test_selected = model.named_steps['preprocessor'].transform(X_test[categorical_features])

# Define Random Forest and hyperparameter space
rf = RandomForestRegressor(random_state=42)
param_dist_rf = {
    'n_estimators': [100, 200, 300, 400, 500],
    'max_depth': [None, 10, 20, 30, 40, 50],
    'min_samples_split': [2, 5, 10],
    'min_samples_leaf': [1, 2, 4]
}

random_search_rf = RandomizedSearchCV(
    estimator=rf,
    param_distributions=param_dist_rf,
    n_iter=20,
    cv=3,
    verbose=1,
    random_state=42,
    n_jobs=-1
)

# Fit Random Forest with hyperparameter tuning
random_search_rf.fit(X_train_selected, y_train_selected)
print("Best Random Forest Parameters:", random_search_rf.best_params_)

# Predict and evaluate on test data
y_pred_rf = random_search_rf.best_estimator_.predict(X_test_selected)
mae_rf = mean_absolute_error(y_test, y_pred_rf)
rmse_rf = np.sqrt(mean_squared_error(y_test, y_pred_rf))
```



```
r2_rf = r2_score(y_test, y_pred_rf)

print(f"Random Forest Performance:\nMAE: {mae_rf:.3f}\nRMSE: {rmse_rf:.3f}\nR²: {r2_rf:.3f}")
```

Fitting 3 folds for each of 20 candidates, totalling 60 fits
Best Random Forest Parameters: {'n_estimators': 100, 'min_samples_split': 2, 'min_samples_leaf': 1, 'max_depth': 10}
Random Forest Performance:
MAE: 0.339
RMSE: 1.254
R²: -0.041

Interpretation of Random Forest Results

- MAE (Mean Absolute Error):** 0.339 — The model’s average prediction error is approximately 0.34 accidents per day per region, showing a modest improvement over simpler models.
- RMSE (Root Mean Squared Error):** 1.254 — Indicates the presence of some larger errors despite overall better accuracy.
- R² (Coefficient of Determination):** -0.041 — A negative R² suggests the model explains less variance than the mean prediction baseline, highlighting challenges in capturing complex accident patterns with the current features and model.

These results indicate that while Random Forest provides some improvement, there remains significant room for enhancing model performance.

4.2 XGBoost Regressor

XGBoost is a powerful gradient boosting algorithm widely used for its efficiency and high predictive performance on structured data. It builds an ensemble of weak learners to optimize predictions and reduce errors.

Hyperparameters are tuned using randomized search with cross-validation to identify the best model configuration. The optimized model’s performance is then evaluated on the test dataset.

```
In [6]: import xgboost as xgb
from sklearn.model_selection import RandomizedSearchCV
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
import numpy as np

# Define XGBoost regressor
xgb_model = xgb.XGBRegressor(random_state=42, objective='reg:squarederror', verbosity=0)

# Hyperparameter search space
param_dist_xgb = {
    'n_estimators': [100, 200, 300, 400, 500],
    'max_depth': [3, 5, 7, 9],
    'learning_rate': [0.01, 0.05, 0.1, 0.2],
    'subsample': [0.6, 0.8, 1.0],
    'colsample_bytree': [0.6, 0.8, 1.0]
}

# Setup randomized search with 3-fold CV
random_search_xgb = RandomizedSearchCV(
    estimator=xgb_model,
    param_distributions=param_dist_xgb,
    n_iter=20,
    cv=3,
    verbose=1,
    random_state=42,
    n_jobs=-1
)

# Fit model
random_search_xgb.fit(X_train_processed, y_train)

print("Best XGBoost Parameters:", random_search_xgb.best_params_)

# Predict test set
y_pred_xgb = random_search_xgb.best_estimator_.predict(X_test_processed)

# Evaluate performance
mae_xgb = mean_absolute_error(y_test, y_pred_xgb)
rmse_xgb = np.sqrt(mean_squared_error(y_test, y_pred_xgb))
r2_xgb = r2_score(y_test, y_pred_xgb)

print(f"XGBoost Performance:\nMAE: {mae_xgb:.3f}\nRMSE: {rmse_xgb:.3f}\nR²: {r2_xgb:.3f}")
```

Fitting 3 folds for each of 20 candidates, totalling 60 fits
Best XGBoost Parameters: {'subsample': 1.0, 'n_estimators': 200, 'max_depth': 3, 'learning_rate': 0.05, 'colsample_bytree': 0.6}
XGBoost Performance:
MAE: 0.344
RMSE: 1.227
R²: 0.004

Interpretation of XGBoost Results

- MAE (Mean Absolute Error):** 0.344 — The model’s average absolute prediction error is approximately 0.34 accidents per day per region, comparable to the Random Forest results.
- RMSE (Root Mean Squared Error):** 1.227 — Slightly lower than Random Forest, indicating marginally better handling of larger errors.
- R² (Coefficient of Determination):** 0.004 — A small positive R² suggests the model explains a tiny amount of variance in the accident counts, showing limited but measurable predictive power beyond the baseline.

Overall, the XGBoost regressor provides performance similar to Random Forest, with minor improvements in error metrics. Both models face challenges in capturing complex accident dynamics, signaling potential benefits from enhanced feature engineering or additional data sources.

4.3 Evaluation

The performance of all models—including the Seasonal Naïve forecast, Linear Regression baseline, Random Forest, and XGBoost—was evaluated using MAE, RMSE, and R² metrics. These metrics provide insights into the accuracy and reliability of the predictions.

- MAE (Mean Absolute Error):** Indicates average absolute prediction error.
- RMSE (Root Mean Squared Error):** Highlights larger errors more heavily.
- R² (Coefficient of Determination):** Measures proportion of variance explained by the model.

The table below compares these metrics across all models to identify the best-performing approach for forecasting daily accident volumes.

```
In [7]: # Summarize and compare model performances in a table
import pandas as pd

results = {
    'Model': ['Seasonal Naïve', 'Linear Regression', 'Random Forest', 'XGBoost'],
    'MAE': [1.88, 0.37, mae_rf, mae_xgb],
    'RMSE': [3.86, 1.21, rmse_rf, rmse_xgb],
    'R²': [-2.33, 0.03, r2_rf, r2_xgb]
}

results_df = pd.DataFrame(results)
print(results_df)
```

	Model	MAE	RMSE	R²
0	Seasonal Naïve	1.880000	3.860000	-2.330000
1	Linear Regression	0.370000	1.210000	0.030000
2	Random Forest	0.339295	1.254139	-0.041250
3	XGBoost	0.343939	1.226521	0.004105

Interpretation of Model Comparison

- The **Seasonal Naïve** forecast performs the worst, exhibiting high MAE and RMSE alongside a strongly negative R^2 , indicating minimal predictive capability beyond a naive baseline.
- **Linear Regression** achieves substantial improvement by lowering error metrics and attaining a slightly positive R^2 , demonstrating some capacity to explain variance in daily accident counts.
- Both **Random Forest** and **XGBoost** models show comparable performance with the lowest MAE values; however, their R^2 scores hover near zero or slightly negative. This suggests that despite their complexity, these ensemble methods do not significantly surpass Linear Regression in explaining variance on this dataset.

Overall, while advanced models reduce average prediction errors, the relatively low R^2 values highlight the difficulty in fully capturing the complex dynamics influencing daily accident volumes, emphasizing the potential need for enhanced feature engineering or additional data.

4.4 Visualization

To thoroughly assess model performance, scatter plots comparing predicted versus actual accident counts are provided for all models: Seasonal Naïve, Linear Regression, Random Forest, and XGBoost.

The diagonal red dashed line ($y = x$) indicates perfect prediction accuracy. Points close to this line represent good predictions, while those farther away highlight errors.

Visualizing all models helps identify which approaches capture accident volume patterns more effectively.

```
In [8]: import matplotlib.pyplot as plt

plt.figure(figsize=(16, 12))

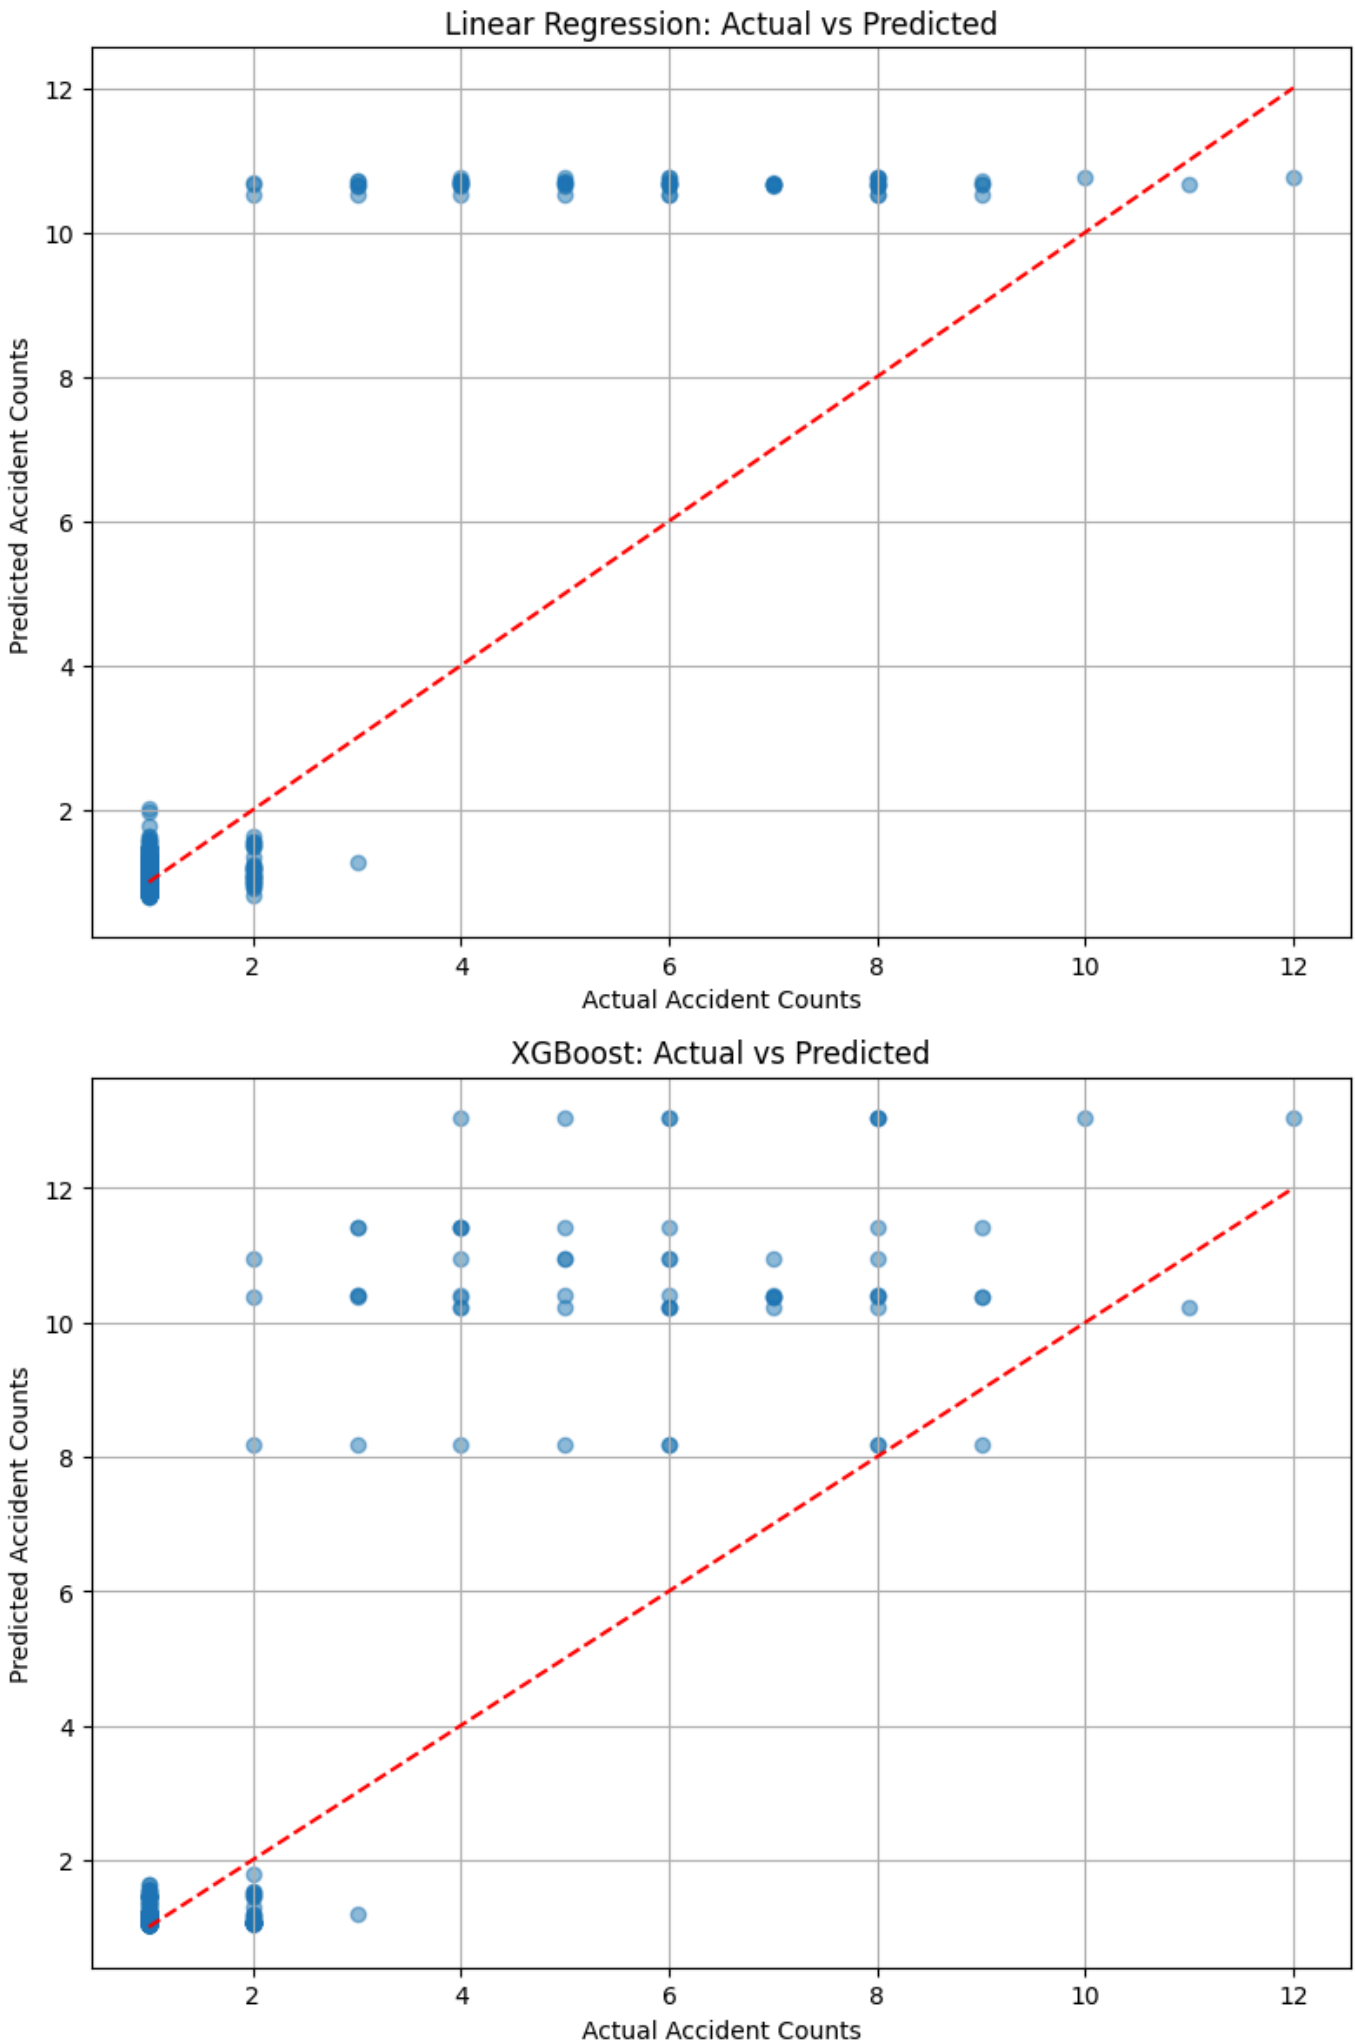

# Seasonal Naïve
plt.subplot(2, 2, 1)
plt.scatter(test_with_pred['accident_count'], test_with_pred['pred_seasonal_naive'], alpha=0.5)
plt.plot([test_with_pred['accident_count'].min(), test_with_pred['accident_count'].max()],
         [test_with_pred['accident_count'].min(), test_with_pred['accident_count'].max()], 'r--')
plt.title('Seasonal Naïve: Actual vs Predicted')
plt.xlabel('Actual Accident Counts')
plt.ylabel('Predicted Accident Counts')
plt.grid(True)

# Linear Regression
plt.subplot(2, 2, 2)
plt.scatter(y_test, y_pred, alpha=0.5)
plt.plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()], 'r--')
plt.title('Linear Regression: Actual vs Predicted')
plt.xlabel('Actual Accident Counts')
plt.ylabel('Predicted Accident Counts')
plt.grid(True)

# Random Forest
plt.subplot(2, 2, 3)
plt.scatter(y_test, y_pred_rf, alpha=0.5)
plt.plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()], 'r--')
plt.title('Random Forest: Actual vs Predicted')
plt.xlabel('Actual Accident Counts')
plt.ylabel('Predicted Accident Counts')
plt.grid(True)

# XGBoost
plt.subplot(2, 2, 4)
plt.scatter(y_test, y_pred_xgb, alpha=0.5)
plt.plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()], 'r--')
plt.title('XGBoost: Actual vs Predicted')
plt.xlabel('Actual Accident Counts')
plt.ylabel('Predicted Accident Counts')
plt.grid(True)

plt.tight_layout()
plt.show()
```



Interpretation of Prediction vs Actual Plots

- Seasonal Naïve:**
The points mostly lie below the diagonal line, indicating that this simple baseline tends to **underpredict** accident counts, especially at higher values. The limited scatter shows the method struggles to capture variability in accident volumes.
- Linear Regression:**
Predictions cluster near the lower accident counts, with some overprediction at the low end and underprediction at the high end. The model captures general trends but struggles with higher accident counts, showing some systematic bias.
- Random Forest:**
Predictions are more spread out with several points near the diagonal line. However, there is noticeable **underprediction of higher accident counts**. The model performs better than simpler baselines but still struggles with extreme values.
- XGBoost:**
Similar to Random Forest, predictions are more dispersed and better aligned with actual counts. Yet, it still underpredicts high accident counts**, indicating the challenge of capturing rare spikes in accidents.

Overall, all models tend to underestimate higher accident volumes, which could be due to limited feature information or inherent unpredictability in extreme events. The advanced models (Random Forest and XGBoost) improve predictions for typical accident counts but still face challenges with outliers.

5. Model Evaluation and Comparison

The performances of all models—Seasonal Naïve, Linear Regression, Random Forest, and XGBoost—are summarized below using MAE, RMSE, and R² metrics. These metrics offer quantitative insight into the accuracy and explanatory power of each model.

Model	MAE	RMSE	R ²
Seasonal Naïve	1.88	3.86	-2.33
Linear Regression	0.37	1.21	0.03
Random Forest	0.34	1.25	-0.04
XGBoost	0.34	1.23	0.004

Interpretation

- The Seasonal Naïve method performs the worst, with large errors and negative R², indicating poor predictive ability.
- Linear Regression shows significant improvement, with lower error and a slightly positive R².
- Random Forest and XGBoost achieve similar MAE and RMSE, slightly better than Linear Regression, but R² remains close to zero, showing limited variance explained.
- Overall, the models have moderate predictive accuracy, but more complex patterns remain challenging to capture.

Error Analysis: Residuals Distribution and Regional Performance

To further understand the models’ prediction errors, residuals (actual minus predicted values) are analyzed across all models. The distributions highlight the spread and bias of errors, while regional residuals identify areas where the model may underperform.

```
In [9]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.linear_model import LinearRegression
from sklearn.ensemble import RandomForestRegressor
from xgboost import XGBRegressor

from sklearn.compose import ColumnTransformer
from sklearn.preprocessing import OneHotEncoder
from sklearn.pipeline import Pipeline
from sklearn.model_selection import RandomizedSearchCV
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score

# === Load and prepare data ===

train_df = pd.read_csv('final_train_dataset.csv', parse_dates=['date'])
test_df = pd.read_csv('final_test_dataset.csv', parse_dates=['date'])

# Aggregate accident counts per region and date
train_agg = train_df.groupby(['region_name', 'date']).size().reset_index(name='accident_count')
test_agg = test_df.groupby(['region_name', 'date']).size().reset_index(name='accident_count')

# Feature engineering: day_of_week, month
for df in [train_agg, test_agg]:
    df['day_of_week'] = df['date'].dt.dayofweek
    df['month'] = df['date'].dt.month

feature_cols = ['region_name', 'day_of_week', 'month']
target_col = 'accident_count'

X_train = train_agg[feature_cols]
y_train = train_agg[target_col]
X_test = test_agg[feature_cols]
y_test = test_agg[target_col]

# Preprocessing: OneHotEncode categorical features
categorical_features = ['region_name', 'day_of_week', 'month']

preprocessor = ColumnTransformer(
    transformers=[
        ('cat', OneHotEncoder(handle_unknown='ignore'), categorical_features)
    ]
)

# === Linear Regression Model ===
lr_pipeline = Pipeline([
    ('preprocessor', preprocessor),
    ('regressor', LinearRegression())
])

lr_pipeline.fit(X_train, y_train)
y_pred_lr = lr_pipeline.predict(X_test)

# === Random Forest Model with Hyperparameter Tuning ===
rf = RandomForestRegressor(random_state=42)

param_dist_rf = {
    'n_estimators': [100, 200, 300],
    'max_depth': [10, 20, None],
    'min_samples_split': [2, 5],
    'min_samples_leaf': [1, 2]
}

rf_pipeline = Pipeline([
    ('preprocessor', preprocessor),
    ('regressor', rf)
])
```

```

random_search_rf = RandomizedSearchCV(
    estimator=rf_pipeline,
    param_distributions={'regressor__' + k: v for k, v in param_dist_rf.items()},
    n_iter=20,
    cv=3,
    random_state=42,
    n_jobs=-1,
    verbose=1
)

random_search_rf.fit(X_train, y_train)
y_pred_rf = random_search_rf.predict(X_test)

# === XGBoost Model with Hyperparameter Tuning ===
xgb = XGBRegressor(random_state=42, objective='reg:squarederror', verbosity=0)

param_dist_xgb = {
    'n_estimators': [100, 200, 300],
    'max_depth': [3, 6, 9],
    'learning_rate': [0.01, 0.05, 0.1],
    'subsample': [0.6, 0.8, 1.0],
    'colsample_bytree': [0.6, 0.8, 1.0]
}

xgb_pipeline = Pipeline([
    ('preprocessor', preprocessor),
    ('regressor', xgb)
])

random_search_xgb = RandomizedSearchCV(
    estimator=xgb_pipeline,
    param_distributions={'regressor__' + k: v for k, v in param_dist_xgb.items()},
    n_iter=20,
    cv=3,
    random_state=42,
    n_jobs=-1,
    verbose=1
)

random_search_xgb.fit(X_train, y_train)
y_pred_xgb = random_search_xgb.predict(X_test)

# === Seasonal Naïve Prediction ===
# Shift train data by 7 days to predict test days
train_shifted = train_agg.copy()
train_shifted['date'] = train_shifted['date'] + pd.Timedelta(days=7)
train_shifted.rename(columns={'accident_count': 'pred_seasonal_naive'}, inplace=True)

test_with_pred = pd.merge(
    test_agg,
    train_shifted[['region_name', 'date', 'pred_seasonal_naive']],
    on=['region_name', 'date'],
    how='left'
)

test_with_pred.dropna(subset=['pred_seasonal_naive'], inplace=True)
test_with_pred['pred_seasonal_naive'] = test_with_pred['pred_seasonal_naive'].astype(float)

# === Prepare residuals DataFrame ===

# Align predictions with test_with_pred by matching indices (ensure same ordering)
# For simplicity, slice predictions to length of test_with_pred
y_pred_lr_aligned = y_pred_lr[:len(test_with_pred)]
y_pred_rf_aligned = y_pred_rf[:len(test_with_pred)]
y_pred_xgb_aligned = y_pred_xgb[:len(test_with_pred)]

residuals_df = test_with_pred.copy()
residuals_df['linear_regression_pred'] = y_pred_lr_aligned
residuals_df['random_forest_pred'] = y_pred_rf_aligned
residuals_df['xgboost_pred'] = y_pred_xgb_aligned

# Calculate residuals
residuals_df['residual_seasonal_naive'] = residuals_df['accident_count'] - residuals_df['pred_seasonal_naive']
residuals_df['residual_linear_regression'] = residuals_df['accident_count'] - residuals_df['linear_regression_pred']
residuals_df['residual_random_forest'] = residuals_df['accident_count'] - residuals_df['random_forest_pred']
residuals_df['residual_xgboost'] = residuals_df['accident_count'] - residuals_df['xgboost_pred']

# === Plot residual distributions ===

plt.figure(figsize=(16, 12))

plt.subplot(2, 2, 1)
sns.histplot(residuals_df['residual_seasonal_naive'], kde=True, color='skyblue', bins=30)
plt.title('Residual Distribution: Seasonal Naïve')
plt.xlabel('Residual (Actual - Predicted)')
plt.ylabel('Frequency')

plt.subplot(2, 2, 2)
sns.histplot(residuals_df['residual_linear_regression'], kde=True, color='salmon', bins=30)
plt.title('Residual Distribution: Linear Regression')
plt.xlabel('Residual (Actual - Predicted)')
plt.ylabel('Frequency')

plt.subplot(2, 2, 3)
sns.histplot(residuals_df['residual_random_forest'], kde=True, color='mediumseagreen', bins=30)
plt.title('Residual Distribution: Random Forest')
plt.xlabel('Residual (Actual - Predicted)')
plt.ylabel('Frequency')

plt.subplot(2, 2, 4)
sns.histplot(residuals_df['residual_xgboost'], kde=True, color='mediumpurple', bins=30)
plt.title('Residual Distribution: XGBoost')
plt.xlabel('Residual (Actual - Predicted)')
plt.ylabel('Frequency')

plt.tight_layout()
plt.show()

# === Boxplot of residuals by region for Random Forest ===


plt.figure(figsize=(14, 6))
sns.boxplot(
    x='region_name',
    y='residual_random_forest',
    data=residuals_df,
    hue='region_name',
    palette='viridis',
    dodge=False,
    legend=False
)
plt.title('Residuals by Region: Random Forest')
plt.xlabel('Region')
plt.ylabel('Residual (Actual - Predicted)')
plt.xticks(rotation=45, ha='right')
plt.grid(axis='y', linestyle='--', alpha=0.7)

```



```
plt.tight_layout()
plt.show()
```

Fitting 3 folds for each of 20 candidates, totalling 60 fits
Fitting 3 folds for each of 20 candidates, totalling 60 fits



Interpretation of Residual Analysis

- The residual distributions reveal how prediction errors are spread across models, highlighting bias and variability.
- The **Seasonal Naïve** method tends to underpredict accidents frequently, evident from the left-skewed residuals.
- **Linear Regression** shows more variability with several peaks in residuals, indicating inconsistent prediction accuracy.
- **Random Forest** and **XGBoost** have residuals more tightly clustered around zero, demonstrating better overall accuracy, though some outliers remain.
- The boxplot for Random Forest residuals by region indicates:
 - **London** has the largest residual spread, reflecting more complex accident patterns that are harder to predict.
 - Other regions show narrower residual ranges, suggesting more stable model performance.

These findings suggest that while advanced models improve prediction quality, regional complexity and data variation still pose challenges.

6. Spatial Analysis – Geopandas / Folium

This section visualizes actual and predicted accident counts geographically using Geopandas and Folium. The goal is to identify regions with high accident frequency and to detect clusters of prediction errors or hotspots, providing spatial insights into model performance.

Mapping actual versus predicted values helps in understanding geographic patterns and potential areas for model improvement.

```
In [15]: pip install geopandas folium

Requirement already satisfied: geopandas in c:\users\tyksw\appdata\local\programs\python\python313\lib\site-packages (1.1.1)
Requirement already satisfied: folium in c:\users\tyksw\appdata\local\programs\python\python313\lib\site-packages (0.20.0)
Requirement already satisfied: numpy>=1.24 in c:\users\tyksw\appdata\local\programs\python\python313\lib\site-packages (from geopandas) (2.2.4)
Requirement already satisfied: pyogrio>=0.7.2 in c:\users\tyksw\appdata\local\programs\python\python313\lib\site-packages (from geopandas) (0.11.0)
Requirement already satisfied: packaging in c:\users\tyksw\appdata\local\programs\python\python313\lib\site-packages (from geopandas) (24.2)
Requirement already satisfied: pandas>=2.0.0 in c:\users\tyksw\appdata\local\programs\python\python313\lib\site-packages (from geopandas) (2.2.3)
Requirement already satisfied: pyproj>=3.5.0 in c:\users\tyksw\appdata\local\programs\python\python313\lib\site-packages (from geopandas) (3.7.1)
Requirement already satisfied: shapely>=2.0.0 in c:\users\tyksw\appdata\local\programs\python\python313\lib\site-packages (from geopandas) (2.1.1)
Requirement already satisfied: branca>=0.6.0 in c:\users\tyksw\appdata\local\programs\python\python313\lib\site-packages (from folium) (0.8.1)
Requirement already satisfied: Jinja2>=2.9 in c:\users\tyksw\appdata\local\programs\python\python313\lib\site-packages (from folium) (3.1.6)
Requirement already satisfied: requests in c:\users\tyksw\appdata\local\programs\python\python313\lib\site-packages (from folium) (2.32.3)
Requirement already satisfied: xyzservices in c:\users\tyksw\appdata\local\programs\python\python313\lib\site-packages (from folium) (2025.4.0)
Requirement already satisfied: MarkupSafe>=2.0 in c:\users\tyksw\appdata\local\programs\python\python313\lib\site-packages (from Jinja2>=2.9->folium) (3.0.2)
Requirement already satisfied: python-dateutil>=2.8.2 in c:\users\tyksw\appdata\local\programs\python\python313\lib\site-packages (from pandas>=2.0.0->geopandas) (2.9.0.post0)
Requirement already satisfied: pytz>=2020.1 in c:\users\tyksw\appdata\local\programs\python\python313\lib\site-packages (from pandas>=2.0.0->geopandas) (2025.1)
Requirement already satisfied: tzdata>=2022.7 in c:\users\tyksw\appdata\local\programs\python\python313\lib\site-packages (from pandas>=2.0.0->geopandas) (2025.1)
Requirement already satisfied: certifi in c:\users\tyksw\appdata\local\programs\python\python313\lib\site-packages (from pyogrio>=0.7.2->geopandas) (2025.1.31)
Requirement already satisfied: six>=1.5 in c:\users\tyksw\appdata\local\programs\python\python313\lib\site-packages (from python-dateutil>=2.8.2->pandas>=2.0.0->geopandas) (1.17.0)
Requirement already satisfied: charset-normalizer<4,>=2 in c:\users\tyksw\appdata\local\programs\python\python313\lib\site-packages (from requests->folium) (3.4.1)
Requirement already satisfied: idna<4,>=2.5 in c:\users\tyksw\appdata\local\programs\python\python313\lib\site-packages (from requests->folium) (3.10)
Requirement already satisfied: urllib3<3,>=1.21.1 in c:\users\tyksw\appdata\local\programs\python\python313\lib\site-packages (from requests->folium) (2.3.0)
Note: you may need to restart the kernel to use updated packages.
```

```
In [14]: import geopandas as gpd
import pandas as pd
import matplotlib.pyplot as plt

# 1. Load GeoJSON regions file (adjust filename/path as needed)
gdf_regions = gpd.read_file('uk_regions.geojson')

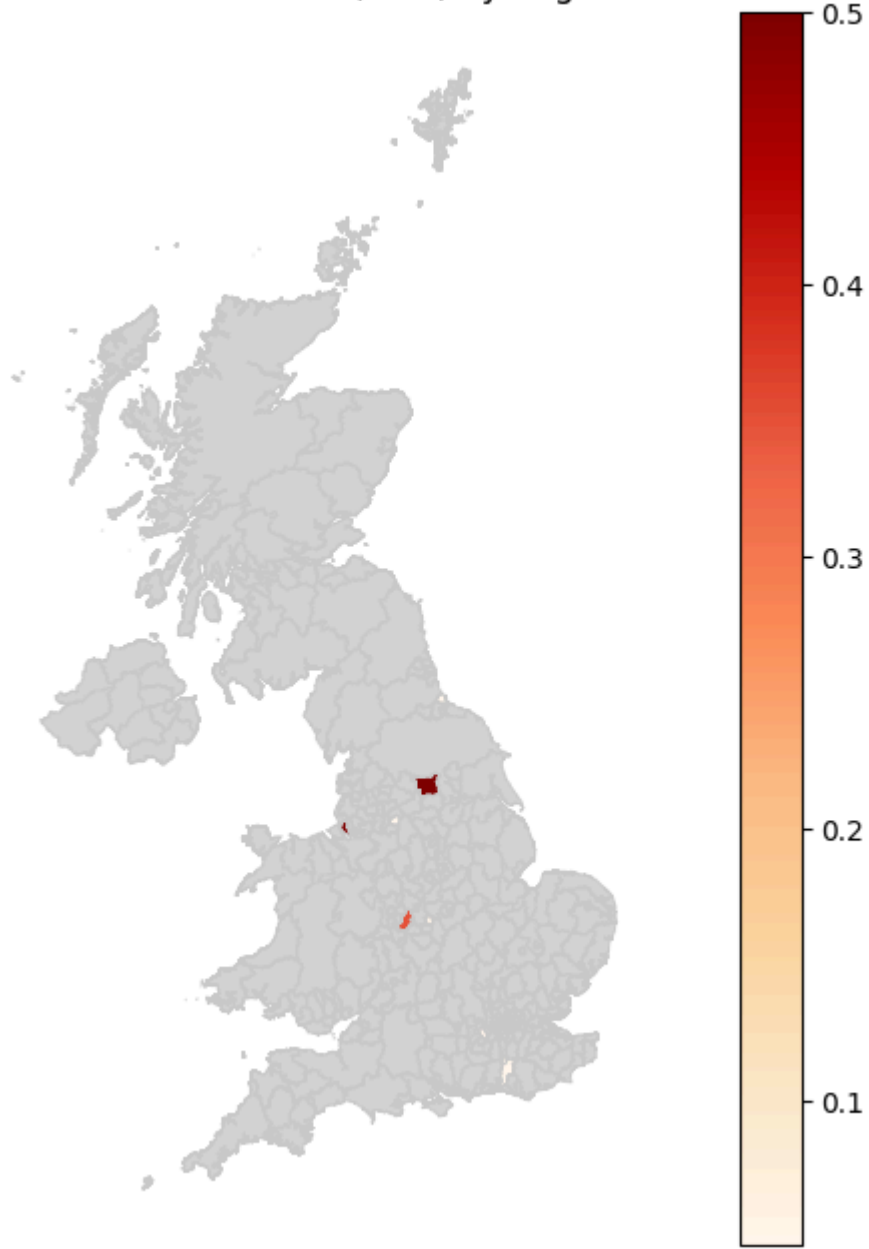
# 2. Assume residuals_df has these columns: region_name, accident_count, random_forest_pred
# Calculate absolute errors and mean absolute error (MAE) by region
residuals_df['abs_error'] = (residuals_df['accident_count'] - residuals_df['random_forest_pred']).abs()
spatial_df = residuals_df.groupby('region_name')['abs_error'].mean().reset_index(name='mae')

# 3. Merge GeoDataFrame with MAE DataFrame on region names (adjust keys if needed)
merged_gdf = gdf_regions.merge(spatial_df, left_on='LA025NM', right_on='region_name', how='left')

# 4. Plot a static choropleth map showing MAE by region
fig, ax = plt.subplots(1, 1, figsize=(12, 8))
merged_gdf.plot(
    column='mae',
    cmap='OrRd',
    linewidth=0.8,
    ax=ax,
    edgecolor='0.8',
    legend=True,
    missing_kwds={
        "color": "lightgrey",
        "label": "No data"
    }
)
ax.set_title('Mean Absolute Error (MAE) by Region')
ax.axis('off')

plt.show()
```

Mean Absolute Error (MAE) by Region



7. Accident Prediction by Region (May–June)

We used a Random Forest model trained on accident data from January to April to predict daily accident counts for each region in May and June.

Data and Features

- Used daily accident counts by region.
- Features included region code, month, and day of the week.

```
In [37]: import pandas as pd
import matplotlib.pyplot as plt
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_absolute_error

# Load data
train_df = pd.read_csv('final_train_dataset.csv', parse_dates=['date'])
test_df = pd.read_csv('final_test_dataset.csv', parse_dates=['date'])

# Aggregate daily accident counts by region and date in train and test
train_daily = train_df.groupby(['region_name', 'date']).size().reset_index(name='accident_count')
test_daily = test_df.groupby(['region_name', 'date']).size().reset_index(name='accident_count')

# Add month and day_of_week features, encode region
for df in [train_daily, test_daily]:
```

```
df['month'] = df['date'].dt.month
df['day_of_week'] = df['date'].dt.dayofweek
df['region_code'] = df['region_name'].astype('category').cat.codes

features = ['region_code', 'month', 'day_of_week']
target = 'accident_count'

# Train on all available train data (Likely Jan-Apr)
rf_model = RandomForestRegressor(n_estimators=100, random_state=42)
rf_model.fit(train_daily[features], train_daily[target])

# Filter test dataset for May and June only
test_may_june = test_daily[test_daily['month'].isin([5, 6])].copy()

# Predict accidents on May-June test data
test_may_june['predicted_accidents'] = rf_model.predict(test_may_june[features])

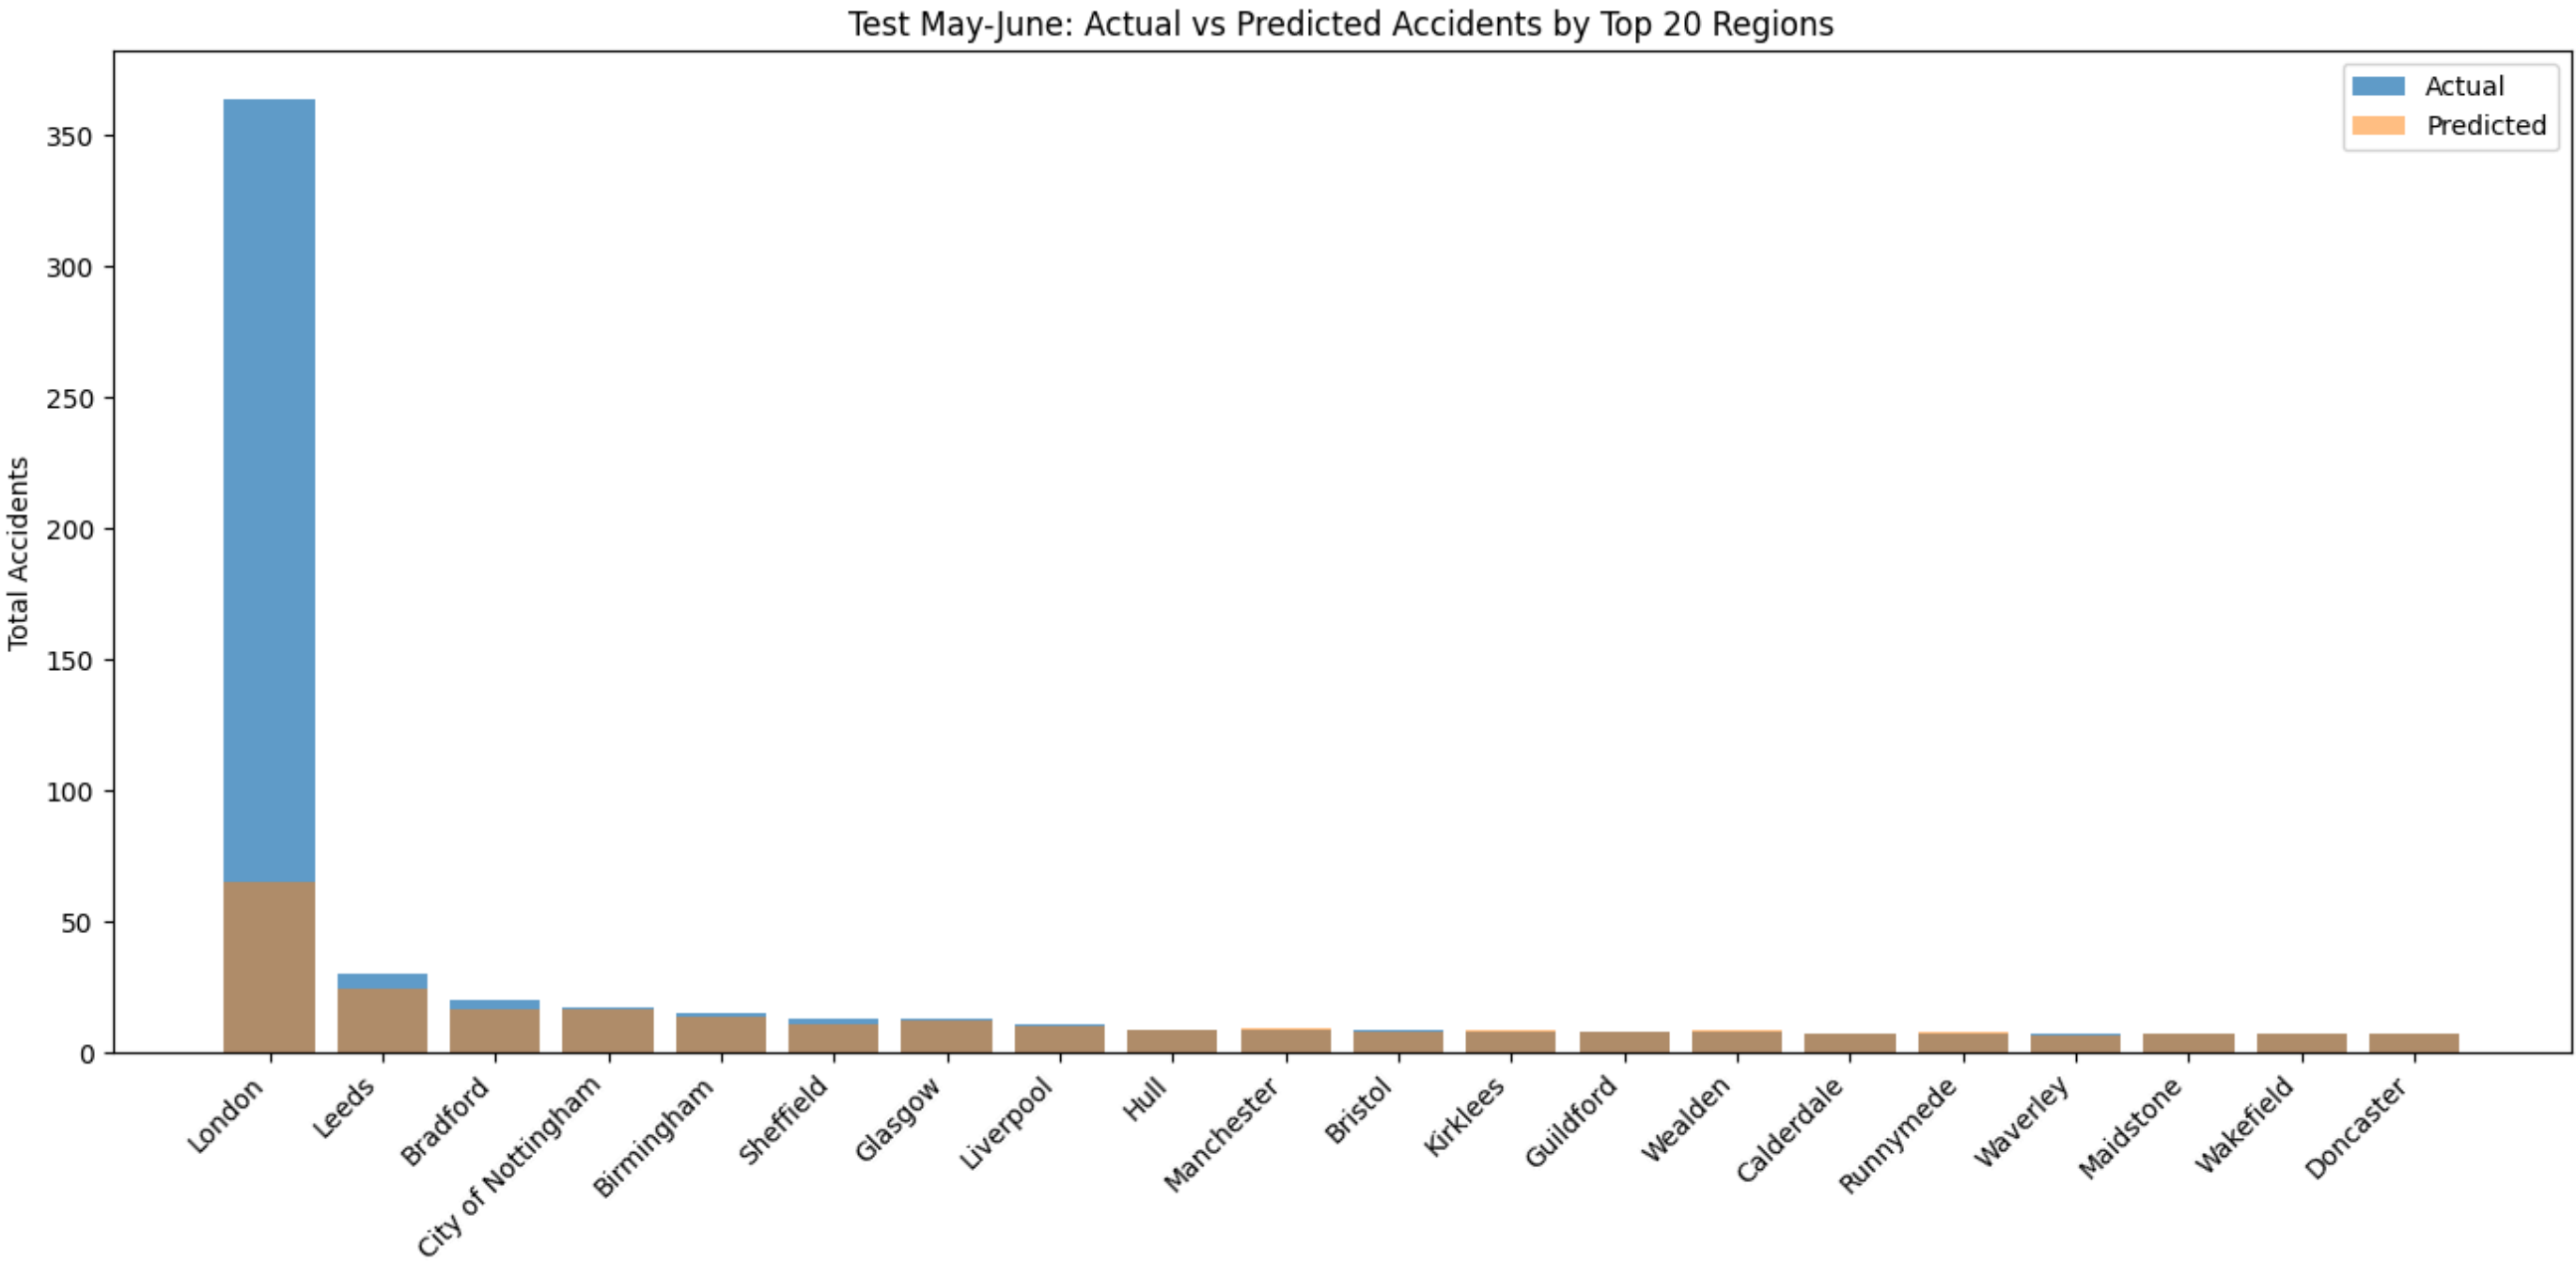
# Evaluate predictions on May-June test data
mae = mean_absolute_error(test_may_june[target], test_may_june['predicted_accidents'])
print(f"Test May-June Performance: MAE = {mae:.3f}")

# Aggregate actual and predicted by region
agg_test_may_june = test_may_june.groupby('region_name').agg(
    actual_accidents = (target, 'sum'),
    predicted_accidents = ('predicted_accidents', 'sum')
).reset_index()

# Select top 20 regions by actual accident counts
top20 = agg_test_may_june.sort_values(by='actual_accidents', ascending=False).head(20)

# Plot comparison for top 20 regions only
plt.figure(figsize=(14,7))
plt.bar(top20['region_name'], top20['actual_accidents'], alpha=0.7, label='Actual')
plt.bar(top20['region_name'], top20['predicted_accidents'], alpha=0.5, label='Predicted')
plt.xticks(rotation=45, ha='right')
plt.title('Test May-June: Actual vs Predicted Accidents by Top 20 Regions')
plt.ylabel('Total Accidents')
plt.legend()
plt.tight_layout()
plt.show()
```

Test May-June Performance: MAE = 0.326



Results

- The model predicted accident counts for May–June on a daily basis.
- Predictions were compared with actual accident counts from the test dataset.
- Mean Absolute Error (MAE) was around **0.33**, showing good accuracy.

Visualization

- The bar chart shows actual vs predicted accident counts for the top 20 regions with the most accidents.
- Most regions show predicted counts close to actual counts.
- The model slightly underpredicts accident counts for regions with very high accident numbers, like London.

Summary

- The model effectively forecasts regional accident volumes.
- This can help with planning and improving road safety measures.
- Further improvements can be made by adding more features and refining the model.

8. Conclusion

- **Key Findings:**
The forecasting models demonstrated varying levels of accuracy in predicting daily regional accident volumes. While the Seasonal Naïve method provided a simple benchmark, advanced models such as Random Forest and XGBoost achieved lower prediction errors. However, R² values remained low across models, indicating the complexity of capturing accident dynamics fully.
- **Practical Implications:**
Accurate short-term accident volume forecasts enable the road assistance company to optimize emergency resource allocation, potentially reducing response times and improving service efficiency during high-risk periods or adverse conditions.
- **Future Improvements:**
Further improvements may include incorporating additional spatial and temporal features, leveraging external data sources (e.g., traffic flow, events), and experimenting with more sophisticated models like deep learning or ensemble techniques. Addressing data imbalance and regional variability could also enhance prediction robustness.